

Problem Set 1: Transportation

Introduction

A colony of super-intelligent alien bioengineers has landed on Earth and has created new species of farm animals. The aliens are performing their experiments on Earth, and plan on transporting the mutant animals back to their home planet. In this problem set, you will implement algorithms to figure out how the aliens should shuttle their experimental animals back across space.

Getting Started

Download ProblemSet1.zip from IVY

Please do not rename the files, change any of the provided helper functions, change function/method names. You will need to keep `ps1_partition.py`, `ps1_sheep_data.txt`, and `ps1_sheep_data_2.txt` in the same folder as `ps1a.py` and `ps1b.py`.

Part A: Transporting Sheep Across Space

The aliens have succeeded in breeding sheep that jump over the moon! Now they want to take home their mutant sheep. The aliens want to take all chosen sheep back, but their spaceship has a weight limit and they want to minimize the number of trips they have to take across the universe. Somehow, the aliens have developed breeding technology to make sheep with only integer weights.

The data for the sheep to be transported is stored in `ps1_sheep_data.txt`, and another set of sheep for another separate transport are in `ps2_sheep_data_2.txt` (There are two files for you to read from and test individually). All of your code for Part A should go into `ps1a.py`.

Problem A.1: Loading Sheep Data [3 Marks]

First, we need to load the sheep data from the data file `ps1_sheep_data.txt`.

The file `ps1_sheep_data_2.txt` is given as another file that you can read and test from (You are also encouraged to write your own sheep list files for testing!), but for now, just work with `ps1_sheep_data.txt`.

You can expect the data to be formatted in pairs of `x,y` on each line, where `x` is the name of the sheep and `y` is a number indicating how much the sheep weighs in tons. Here are the first few lines of `ps1_sheep_data.txt`:

Maggie,3

Herman,7

Betsy,9

...

You can assume that all the sheep have unique names.

Implement the function `load_sheep(filename)` in `ps1a.py`. It should take in the name of the data text file as a string, read in its contents, and return a dictionary that maps sheep names to their weights.

Problem A.2: Greedy Sheep Transport [4 Marks]

One way of transporting sheep is to always pick the heaviest sheep that will fit onto the spaceship first. This is an example of a greedy algorithm. So if there is only 2 tons of free space on your spaceship, with one sheep that is 3 tons and another that's 1 ton, the 1 ton sheep will still get put onto the spaceship.

Implement the greedy algorithm for transporting the sheep back across space in `greedy_sheep_transport`. The result should be a list of lists, where each inner list represents a trip and contains the names of sheep taken on that trip.

Note:

- Make sure not to mutate the dictionary of sheep that is passed in.

Assumptions:

- The order of the list of trips does not matter. That is, `[[1,2],[3,4]]` and `[[3,4],[1,2]]` are considered equivalent list of trips.
- All the sheep are between 0 and 10 tons in weight
- All the sheep have unique names
- If multiple sheep weigh the same amount, break ties by name
- the spaceship has a cargo weight limit (in tons), which is passed into the function as a parameter

Example:

Suppose the spaceship has a weight limit of 10 tons and the set of sheep to transport is `{"Jesse":6, "Maybel":3, "Callie":2, "Maggie":5}`.

The greedy algorithm will first pick Jesse as the heaviest sheep for the first trip. There is still space for 4 tons on the trip. Since Maggie will not fit on this trip, the greedy algorithm picks Maybel, the heaviest sheep that will still fit. Now there is only 1 ton of space left, and none of the sheep can fit in that space, so the first trip is `["Jesse", "Maybel"]`.

For the second trip, the greedy algorithm first picks Maggie as the heaviest remaining sheep, and then picks Callie as the last sheep. Since they will both fit, this makes the second trip `["Maggie", "Callie"]`.

The final result then is `[["Jesse", "Maybel"], ["Maggie", "Callie"]]`.

Problem A.3: Brute Force Sheep Transport [5 Marks]

Another way to transport sheep is to look at every possible combination of trips and pick the best one. This is an example of a brute force algorithm.

Implement a brute force algorithm to find the minimum number of trips needed to take all the sheep across the universe in `brute_force_sheep_transport`. The result should be a list of lists, where each inner list represents a trip and contains the names of sheep taken on that trip.

Notes:

- Make sure not to mutate the dictionary of sheep.
- In order to enumerate all possible combinations of trips, you will want to work with set partitions. You are provided with a helper function called `get_partitions` that generates all the set partitions for a set of sheep. More details on this function are provided below.

Assumptions:

- Again, you can assume that order does not matter. $[[1,2],[3,4]]$ and $[[3,4],[1,2]]$ are considered equivalent lists of trips, and $[[1,2],[3,4]]$ and $[[2,1],[3,4]]$ are considered the same partitions of $[1,2,3,4]$.
- You can assume that all the sheep are between 0 and 10 tons in weight
- All the sheep have unique names
- If multiple sheep weigh the same amount, break ties by name
- The spaceship has a cargo weight limit (in tons), which is passed into the function as a parameter

Helper function `get_partitions`:

To generate all the possibilities for the brute force method, you will want to work with set partitions. For instance, all the possible 2-partitions of the list $[1,2,3,4]$ are $[[1,2],[3,4]]$, $[[1,3],[2,4]]$, $[[2,3],[1,4]]$, $[[1],[2,3,4]]$, $[[2],[1,3,4]]$, $[[3],[1,2,4]]$, $[[4],[1,2,3]]$.

To help you with creating partitions, a helper function `get_partitions(list)` is provided. It takes as input a list and returns a generator that contains all the possible partitions of this list, from 0 partitions to n -partitions where n is the length of this list.

To use generators, you must iterate over the generator to retrieve the elements; you cannot index into a generator. For instance, the recommended way to call `get_partitions` on a list $[1,2,3]$ is the following:

```
for partition in get_partition([1,2,3]):  
    print(partition)
```

Try out this snippet of code to see what is printed.

Generators are outside the scope of Y4CEP.

Example:

Suppose the spaceship has a cargo weight limit of 10 tons and the set of sheep to transport is $\{ \text{"Jesse": 6, "Maybel": 3, "Callie": 2, "Maggie": 5} \}$.

The brute force algorithm will first try to fit them on only one trip, $[\text{"Jesse"}, \text{"Maybel"}, \text{"Callie"}, \text{"Maggie"}]$. Since this trip contains 16 tons of sheep, it is over the weight limit and does not work.

Then the algorithm will try fitting them on all combinations of two trips. Suppose it first tries $[[\text{"Jesse"}, \text{"Maggie"}], [\text{"Maybel"}, \text{"Callie"}]]$. This solution will be rejected because Jesse and Maggie together are over the weight limit and cannot be on the same trip. The algorithm will continue trying two trip partitions until it finds one that works, such as $[[\text{"Jesse"}, \text{"Callie"}], [\text{"Maybel"}, \text{"Maggie"}]]$.

The final result is then $[[\text{"Jesse"}, \text{"Callie"}], [\text{"Maybel"}, \text{"Maggie"}]]$.

Note that depending on which sheep it tries first, the algorithm may find a different, optimal solution. Another optimal result could be $[[\text{"Jesse"}, \text{"Maybel"}], [\text{"Callie"}, \text{"Maggie"}]]$.

Problem A.4: Comparing the Sheep Transport Algorithms [3 Marks]

Implement `compare_sheep_transport_algorithms`. Load the sheep data in `ps1_sheep_data.txt`, and then run your greedy and brute force sheep transport algorithms on the data to find the minimum number of trips found by each algorithm and how long each method takes. Use the default weight limits of 10 for both algorithms.

Note: Make sure you have tested both your greedy and brute force algorithms before implement this.

Hints:

- you can measure the rime a block of code takes to execute using the `time.time()` function as follows:

```
start = time.time()
## code to be timed
end = time.time()
print end - start
```

This will print the duration in seconds, as a float.
- using the given default weight limits of 10 and the given sheep data, both algorithms should not take more than a few seconds to run.

Problem A.5: Writeup [3 Marks]

Answer the following questions in a PDF file called `ps1_answers.pdf`.

1. What were your results from `compare_sheep_transport_algorithms`? Which algorithm runs faster? Why?
2. Does the greedy algorithm return the optimal solution? Why/why not?
3. Does the brute force algorithm return the optimal solution? Why/why not?

Part B: Silver Eggs

After the aliens transport the sheep, one of their interns finds flock of silver ducks. Due to budget cuts they are forced to downsize their ships so they cannot simply take the ducks back, but instead decide to take their silver eggs back. Their ships can only hold a certain amount of weight, and are very small inside. So, because all the eggs are the same size, but have different weights, they want to bring back as few eggs as possible that fill their ship's weight limit. Silver eggs are all the same size, but may have different densities, thus 1 two pound egg is better than 2 one-pound eggs.

Problem B.1: Dynamic Programming: The Plan [8 marks]

The aliens want to carry as few eggs as possible on their trip as they do not have a lot of space on their ships. The have taken detailed notes on the weights of all the eggs that ducks can lay in a given flock and how much weight their ships can hold.

Implement a dynamic programming algorithm to find the minimum number of eggs needed to make a given weight for a certain ship in `dp_make_weight`. The result should be an integer representing the minimum number of eggs from the given flock of ducks needed to make the given weight. You algorithm does not need to return what the weight of the eggs are, just the minimum number of eggs.

A solution that does not use dynamic programming will receive zero points for this problem.

Notes:

- If you try using a brute force algorithm on this problem, it will take a substantially long time to generate the correct output if there are a large number of egg weight available.
- you may implement your algorithm using the top-down recursive method or the bottom-up tabulation method. The former was covered in lesson, but we will accept either method.
- The memo parameter in `dp_make_weight` is optional. You may or may not need to use this parameter depending on your implementation.

Assumptions:

- All the eggs weights are unique between different ducks, but a given duck will always lay the same size egg
- the aliens can wait around for the ducks to lay as many eggs as they need.

Example:

Suppose the first ship can carry 99 pounds and uses the first flock of ducks they find, which contains ducks that lay eggs of weights 1, 5, 10 and 20 pounds.

Your dynamic programming algorithm should return 10 (the minimum number of egg needed to make 99 pounds is 4 eggs of 20 pounds, 1 egg of 10 pounds, 1 egg of 5 pounds, and 4 eggs of 1 pound).

Hints:

- Dynamic programming involves breaking a larger problem into smaller simpler subproblems, solving the subproblems, and storing their solutions. What are the subproblems in this case? What values do you want to store?
- This problem is analogous to the knapsack problem. Imagine the eggs are items you are packing. What is the objective function? What is the weight limit in this case? What are the values of each item? What is the weight of each item?

Problem B.2: Writeup [4 Marks]

Answer the following questions in the same PDF file `ps1_answers.pdf`.

1. Explain why it would be difficult to use a brute force algorithm to solve this problem if there were 30 different egg weights. You do not need to implement a brute force algorithm in order to answer this.
2. If you were to implement a greedy algorithm for finding the minimum number of eggs needed, what would the objective function be? What would the constraints be? What strategy would your greedy algorithm follow to pick which eggs to take? You do not need to implement a greedy algorithm in order to answer this.
3. Will a greedy algorithm always return the optimal solution to this problem? Explain why it is optimal or give an example of when it will not return the optimal solution. Again, you do not need to implement a greedy algorithm in order to answer this.